

man bpf

The `bpf()` system call performs a range of operations related to extended Berkeley Packet Filters. Extended BPF (or eBPF) is similar to the original ("classic") BPF (cBPF) used to filter network packets.

For both cBPF and eBPF programs, the kernel statically analyzes the programs before loading them, in order to ensure that they cannot harm the running system.

eBPF extends cBPF in multiple ways, including the ability to call a fixed set of in-kernel helper functions and access shared data structures such as eBPF maps.



eBPF

“Superpowers have finally come to Linux”

“eBPF does to Linux what JavaScript does to HTML”

- Brendan Gregg, Netflix





eBPF

Run code in the kernel
without having to write a kernel module



man bpf

eBPF programs can be written in a **restricted C** that is compiled (using the **clang** compiler) into **eBPF bytecode**. Various features are omitted from this restricted C, such as loops, global variables, variadic functions, floating-point numbers, and passing structures as function arguments.



man bpf

The kernel contains a just-in-time (JIT) compiler that translates eBPF bytecode into native machine code for better performance.





Writing hello world

Avoiding tool chain hell

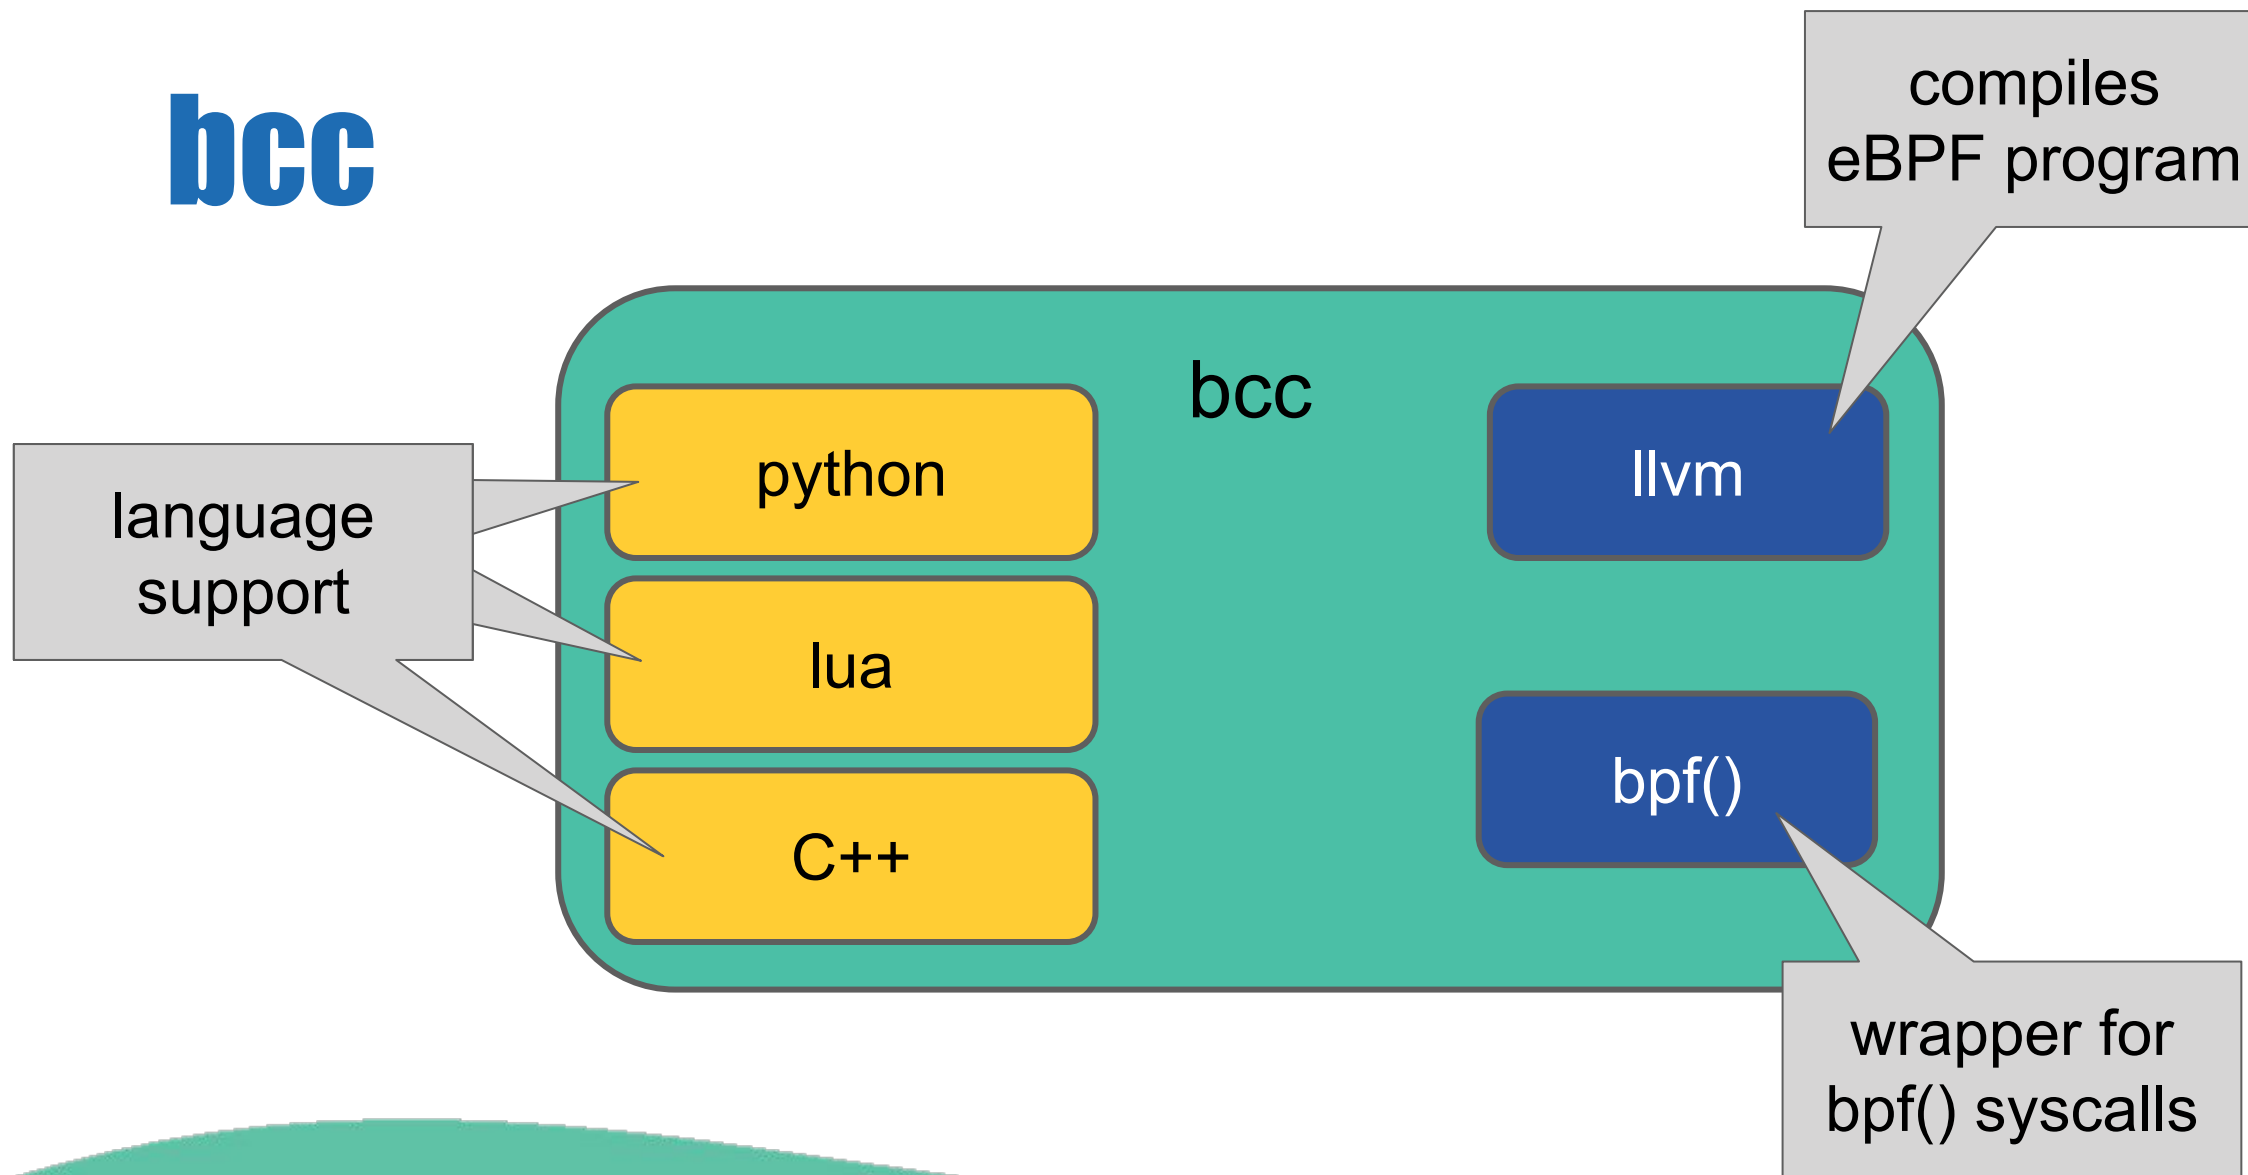
@lizrice

bcc

“BCC **makes BPF programs easier to write**, with kernel instrumentation in C (and includes a C wrapper around LLVM), and front-ends in Python and lua.”



bcc




```
#!/usr/bin/python
from bcc import BPF

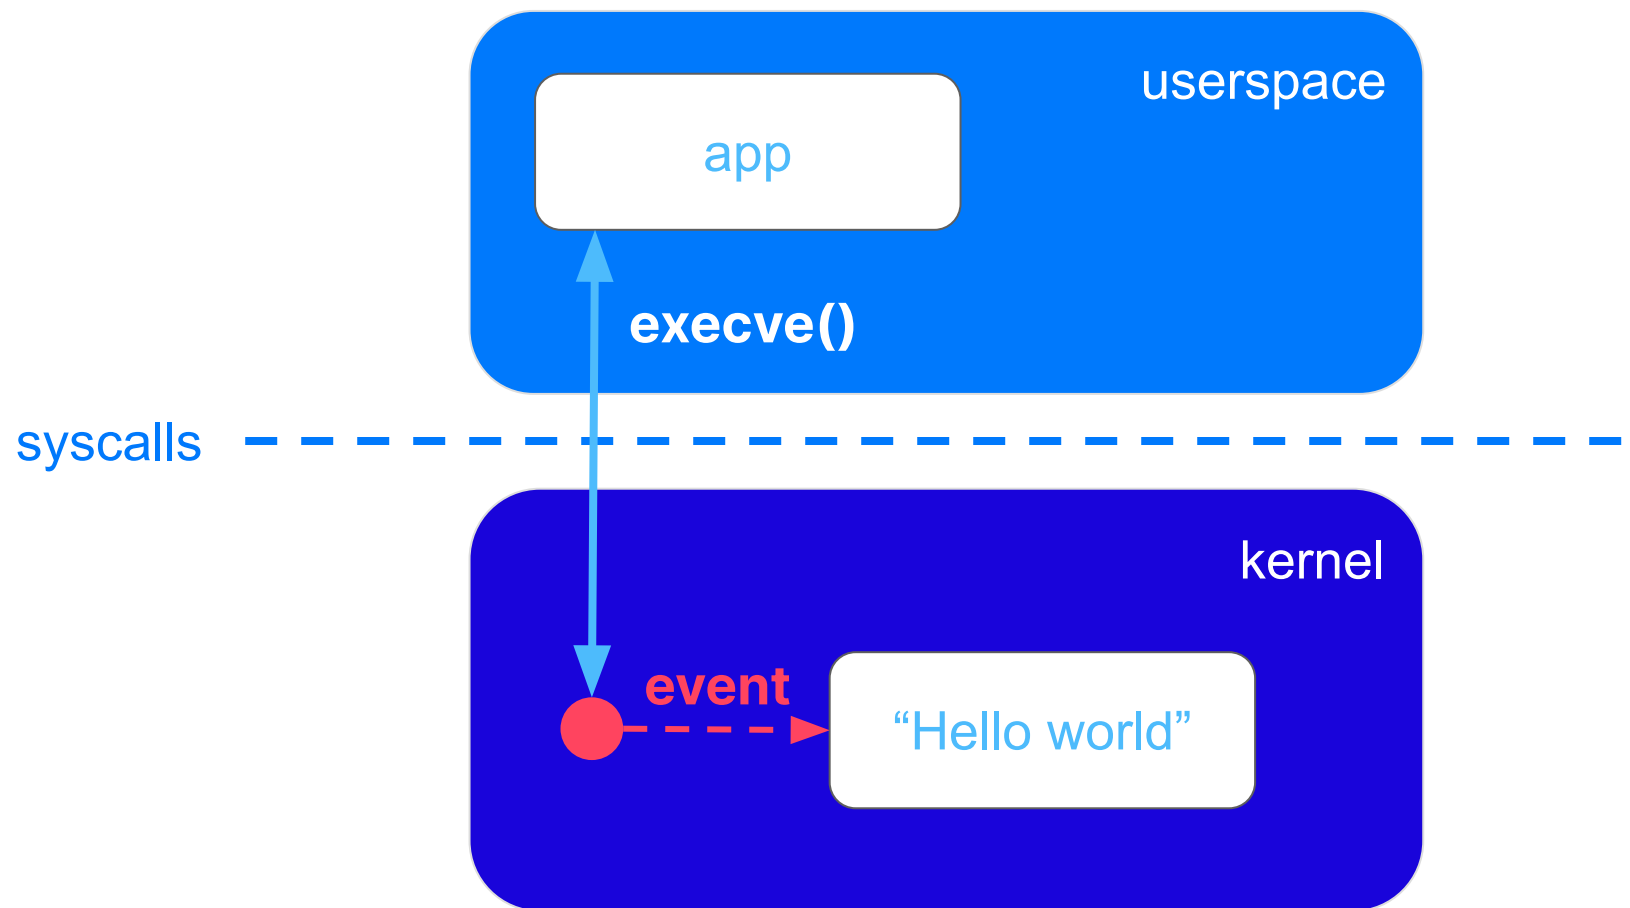
prog = """
int my_prog(void *ctx) {
    bpf_trace_printk("Hello world\\n");
    return 0;
}
"""

b = BPF(text=prog)
b.attach_kprobe(event="sys_clone", fn_name="my_prog")
b.trace_print()
```

Use strace to see
the system calls



eBPF Hello World





Attaching BPF programs to events

@lizrice

Triggering eBPF programs

eBPF programs can be attached to different events.

- Kprobes
- Uprobes
- Tracepoints
- Network packets
- Perf events
- etc...




bpf_trace_printk0

Writing to

```
/sys/kernel/debug/tracing/trace_pipe
```



@lizrice



COALITION FOR THE

eBPF helper functions

These helpers are used by eBPF programs to **interact with the system, or with the context in which they work**. For instance, they can be used to print debugging messages, to get the time since the system was booted, to interact with eBPF maps, or to manipulate network packets.

```
bpf_trace_printk()  
bpf_map_*_elem()  
bpf_get_current_pid_tgid()  
...
```

